

EKS Cluster Creation with Terraform – Full Production Guide (Code Included)

1. Purpose

This document provides a **complete, production-ready EKS setup using Terraform**, including:

- EKS cluster creation
- Networking & IAM
- IRSA (IAM Roles for Service Accounts)
- AWS Load Balancer Controller
- GitOps with Argo CD
- Multi-region DR
- FinOps cost optimization

This is designed for **real-world environments**, not demos.

2. Prerequisites

- AWS account
 - AWS CLI configured
 - Terraform $\geq 1.x$
 - kubectl
 - Helm
-

3. Terraform Project Structure

```
eks-terraform/
├── backend.tf
├── provider.tf
├── variables.tf
├── vpc.tf
├── iam.tf
├── eks.tf
├── nodegroup.tf
├── irsa.tf
├── autoscaler.tf
└── externaldns.tf
```

```
└── secrets.tf  
└── outputs.tf
```

4. Terraform Backend (Remote State)

```
terraform {  
  backend "s3" {  
    bucket      = "eks-terraform-state"  
    key         = "eks/terraform.tfstate"  
    region      = "us-east-1"  
    dynamodb_table = "terraform-locks"  
    encrypt     = true  
  }  
}
```

5. Provider Configuration

```
provider "aws" {  
  region = var.region  
  default_tags {  
    tags = {  
      Environment = "production"  
      Owner      = "platform"  
      CostCenter = "eks"  
    }  
  }  
}
```

6. VPC and Networking (vpc.tf)

```
module "vpc" {  
  source  = "terraform-aws-modules/vpc/aws"  
  name    = "eks-vpc"  
  cidr   = "10.0.0.0/16"  
  
  azs      = ["us-east-1a", "us-east-1b"]  
  private_subnets = ["10.0.1.0/24", "10.0.2.0/24"]  
  public_subnets = ["10.0.101.0/24", "10.0.102.0/24"]  
  
  enable_nat_gateway = true  
  single_nat_gateway = true
```

```
}
```

Best Practice: **Worker nodes always in private subnets.**

7. IAM Roles (iam.tf)

EKS Cluster Role

```
resource "aws_iam_role" "eks_cluster_role" {
  name = "eks-cluster-role"
  assume_role_policy = jsonencode({
    Version = "2012-10-17"
    Statement = [
      {
        Effect = "Allow"
        Principal = { Service = "eks.amazonaws.com" }
        Action = "sts:AssumeRole"
      }
    ]
  })
}
```

Attach policies:

- AmazonEKSClusterPolicy
 - AmazonEKSVPCResourceController
-

8. EKS Cluster (eks.tf)

```
module "eks" {
  source  = "terraform-aws-modules/eks/aws"
  cluster_name  = "prod-eks"
  cluster_version = "1.29"

  vpc_id    = module.vpc.vpc_id
  subnet_ids = module.vpc.private_subnets

  enable_irsa = true
}
```

9. Managed Node Groups (nodegroup.tf)

```
module "eks" {
  eks_managed_node_groups = {
    default = {
      instance_types = ["t3.medium"]
      desired_size  = 2
      max_size      = 5
      min_size      = 1
      capacity_type = "ON_DEMAND"
    }
    spot = {
      instance_types = ["t3.large"]
      desired_size  = 1
      max_size      = 3
      min_size      = 0
      capacity_type = "SPOT"
    }
  }
}
```

10. Configure kubectl

```
aws eks update-kubeconfig --region us-east-1 --name prod-eks
kubectl get nodes
```

11. IRSA – IAM Roles for Service Accounts (irsa.tf)

```
module "alb_irsa" {
  source = "terraform-aws-modules/iam/aws//modules/iam-role-for-service-accounts-eks"

  role_name = "alb-controller"

  attach_load_balancer_controller_policy = true

  oidc_providers = {
    main = {
      provider_arn      = module.eks.oidc_provider_arn
      namespace_service_accounts = ["kube-system:aws-load-balancer-controller"]
    }
  }
}
```

12. AWS Load Balancer Controller (Helm)

```
helm repo add eks https://aws.github.io/eks-charts
helm install aws-load-balancer-controller eks/aws-load-balancer-controller \
-n kube-system \
--set clusterName=prod-eks \
--set serviceAccount.create=false \
--set serviceAccount.name=aws-load-balancer-controller
```

13. GitOps with Argo CD

Install Argo CD

```
kubectl create namespace argocd
kubectl apply -n argocd \
-f https://raw.githubusercontent.com/argoproj/argo-cd/stable/manifests/install.yaml
```

14. GitOps Repository Structure

```
gitops-repo/
├── apps/
│   ├── frontend/
│   ├── backend/
│   └── database/
└── argocd-apps/
```

15. Argo CD Application Manifest

```
apiVersion: argoproj.io/v1alpha1
kind: Application
metadata:
  name: frontend
  namespace: argocd
spec:
  project: default
  source:
    repoURL: https://github.com/org/gitops-repo
    path: apps/frontend
    targetRevision: main
  destination:
    server: https://kubernetes.default.svc
```

```
namespace: frontend
syncPolicy:
  automated:
    prune: true
    selfHeal: true
```

16. Production Best Practices

- Private API endpoint for EKS
 - IRSA for all controllers
 - Image scanning before deploy
 - Resource limits on pods
 - HPA + Cluster Autoscaler
 - Multi-region DR planning
 - FinOps cost controls
-

17. Multi-Region Disaster Recovery (DR)

Strategy

- Active-Passive EKS clusters in two regions
- Route 53 health-check-based failover
- GitOps keeps clusters synchronized automatically

Primary: `us-east-1`

Secondary: `us-west-2`

Route 53 Failover

```
resource "aws_route53_health_check" "primary" {
  fqdn      = "app.example.com"
  port      = 443
  type      = "HTTPS"
  resource_path = "/health"
  failure_threshold = 3
  request_interval = 30
}
```

```
resource "aws_route53_record" "primary" {
  zone_id = var.zone_id
  name   = "app.example.com"
  type   = "A"
  set_identifier = "primary"
```

```

failover_routing_policy { type = "PRIMARY" }
alias { name = aws_lb.primary.dns_name; zone_id = aws_lb.primary.zone_id;
evaluate_target_health = true }
}

resource "aws_route53_record" "secondary" {
zone_id = var.zone_id
name   = "app.example.com"
type   = "A"
set_identifier = "secondary"
failover_routing_policy { type = "SECONDARY" }
alias { name = aws_lb.secondary.dns_name; zone_id = aws_lb.secondary.zone_id;
evaluate_target_health = true }
}

```

Argo CD Multi-Cluster

argocd cluster add primary-context
argocd cluster add secondary-context

ApplicationSet for multi-region deployments ensures manifests are synced in both clusters automatically.

18. FinOps Cost Controls

Node Groups & Scaling

- Use **mixed instance types**
- Use **Spot instances** for non-critical workloads
- HPA + Cluster Autoscaler reduces waste

capacity_type = "SPOT"

Network & Storage Costs

- Prefer **S3 over EFS**
- Monitor NAT Gateway traffic
- Use VPC endpoints where possible

CI/CD Costs

- Enforce image size limits
- Remove unused images from ECR

- Scan images before build

Observability for FinOps

- Prometheus/Grafana dashboards for utilization and idle resources
 - Track underutilized workloads and over-provisioned nodes
-

19. Final Architecture Outcome

- Multi-region DR with automatic failover
 - GitOps sync across clusters
 - Autoscaling at pod and node layers
 - Secure secrets via Secrets Manager + CSI Driver
 - ALB ingress with ExternalDNS
 - Full FinOps visibility and cost optimization
-